

Introduction to R for Data Analysis (Hands-on Workshop)



M.Sc. Vijay Palliyil





TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt



M. Sc. Vijay Palliyil
palliyil@verkehr.tu-darmstadt.de

Learning Goals

- #1 Running R code
- #2 Packages & data import
- #3 Wrangling data
- #4 Visualising insights
- #5 How to keep learning R and advance



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt

Lecture Structure



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt

- #1 Orientation
- #2 Rstudio & R basics
- #3 Import and inspect data
- #4 Wrangling with dplyr
- #5 Visualisation with ggplot2
- #6 Exporting, debugging and way forward

#1 Orientation



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt

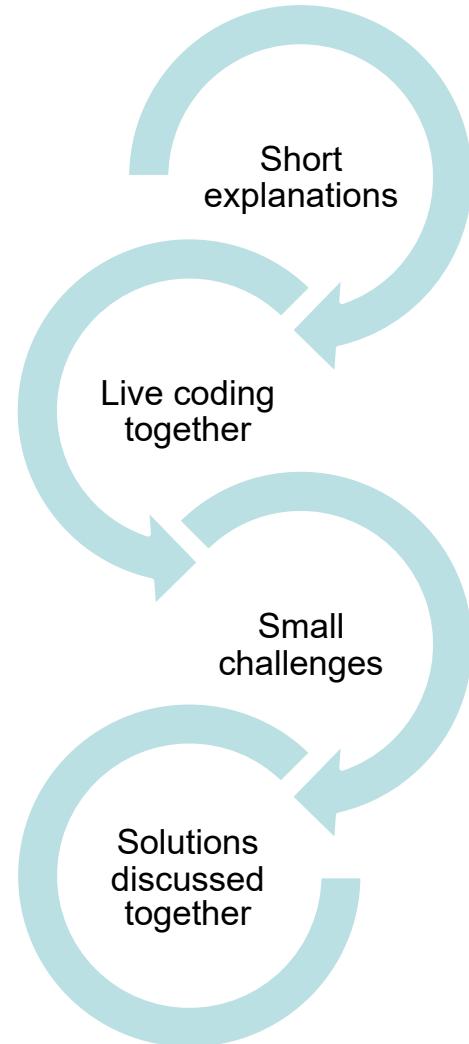
“R is a language for data analysis and graphics.”

- *Ross Ihaka, co-creator of R*

“Friends don’t let friends analyze data in Excel. They let them do it in R.”

- *A wise data analyst*

How this session will work:



Reasons to learn R

- Free and open source.
- Software for data science:
 - experiment/survey design
 - data retrieval
 - data wrangling
 - data analysis
 - reporting
- A programming language, so we can use existing functions to code up our data science tasks or write new functions for customised/novel tasks

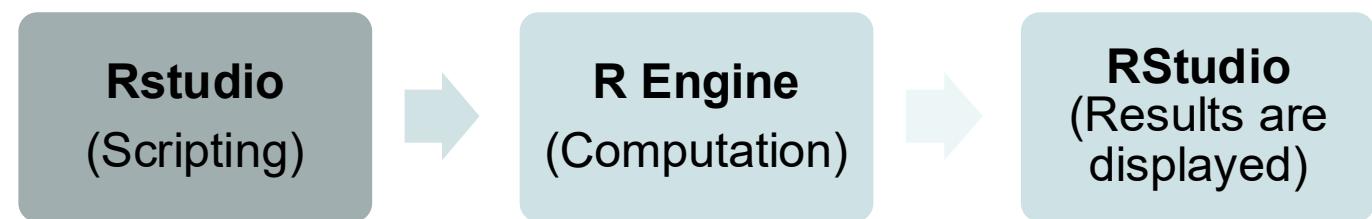


Tools we need installed:

R and Rstudio

What is R?

- A programming language
- A statistical computing engine
- Executes your code
- Produces results



What is Rstudio?

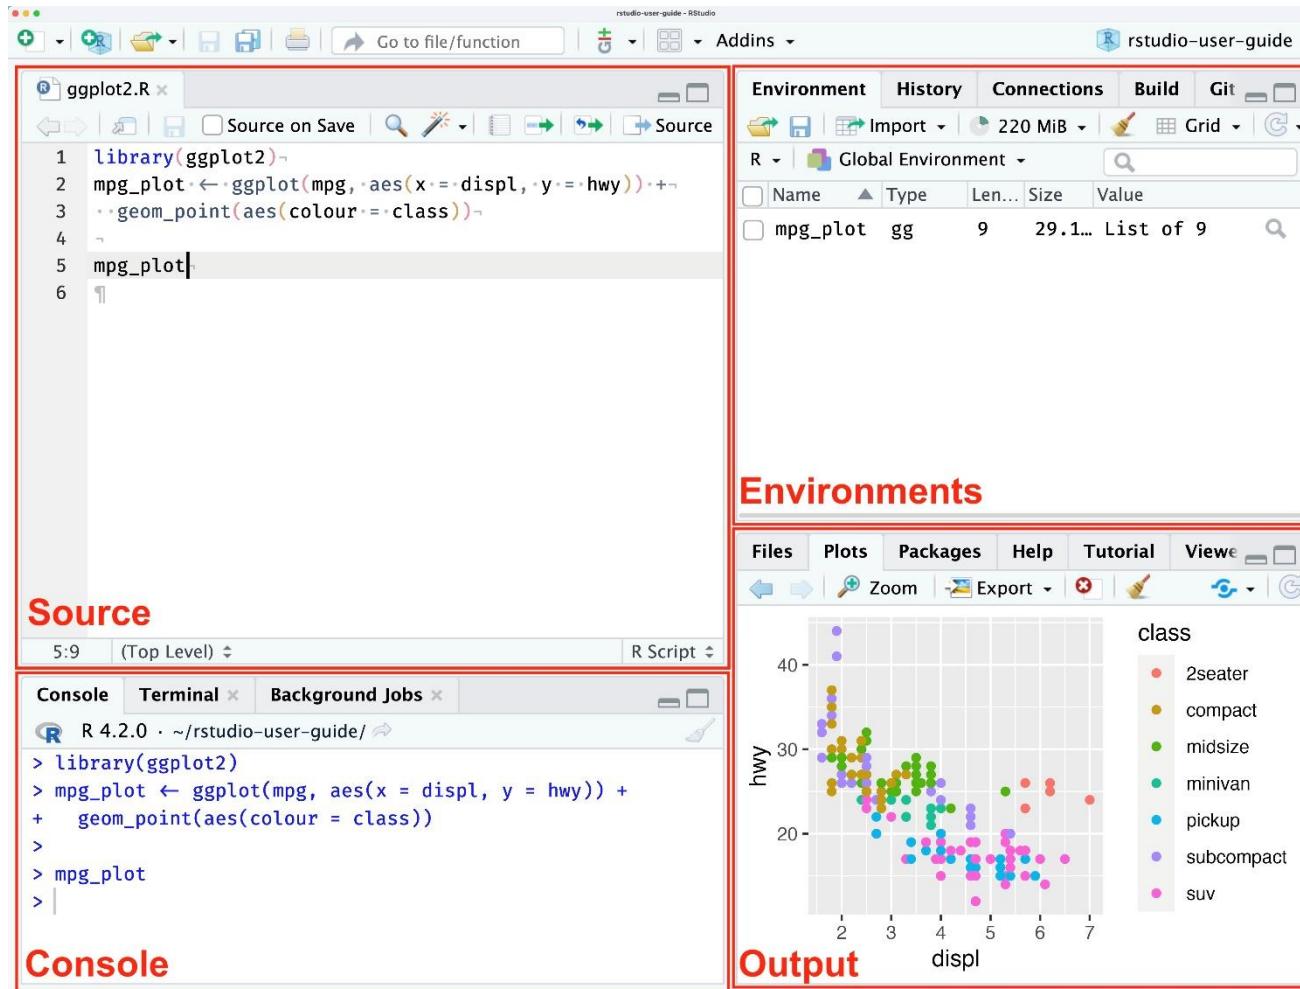
- An IDE for R
- Makes R easier to use
- Editor, console, plots, help
- Does NOT replace R



#2 RStudio & R basics

RStudio IDE (Integrated Development Environment):

The **Source pane** is where you can edit and save R or Python scripts or author computational documents.



The **Console pane** is used to write short interactive R commands.

The **Environment pane** displays temporary R objects as created during that R session.

The **Output pane** displays the plots, tables, or HTML outputs of executed code along with files saved to disk.

Image source: <https://docs.posit.co/ide/user/ide/get-started/>



TECHNISCHE
UNIVERSITÄT
DARMSTADT

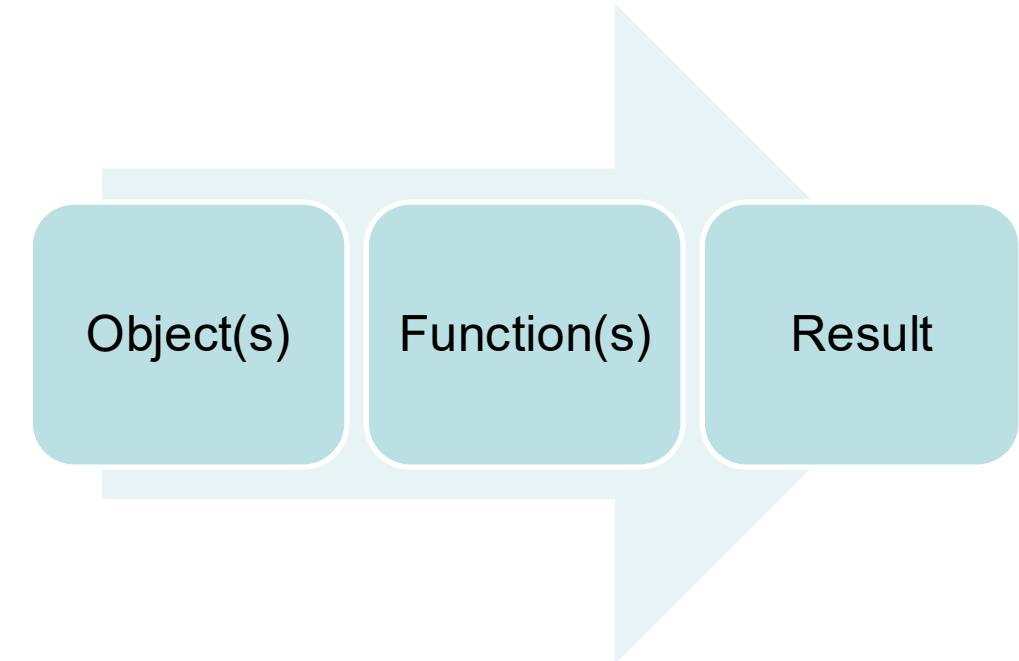


Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt

How R works (conceptually):

R and Rstudio

- You create **objects** (values, variables, tables, etc.)
- Objects have **names**
- **Functions** take objects as input
- Functions return new **objects or results**
- Functions in R are written to operate on whole vectors at once
- A **vector** in R is an ordered collection of values of the same type. A vector is R's way of storing many values under one name.
- Columns in data tables are vectors



First R code-along

- Create an object with the assign symbol <-
 - <- assigns
- Use functions with ()
 - Example: **mean()**
- Run code line by line
 - Type into the script pane
 - **Ctrl + Enter** or **Cmd + Enter** to run the line
- R prints or returns results
- **R is case sensitive!**

The screenshot shows the RStudio interface. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, and Help. Below the menu is a toolbar with various icons. The main area has a script pane titled "Untitled1" containing the following R code:

```
1 # The line below creates an object x -- a vector of numbers
2
3 x <- c(1,5,2,10)
4
5 # The next line applies the mean function to that vector
6
7 mean(x)
8
9 # next line shows other possible calculations
10
11 x^2
12 y <- x^2
13 y
14
```

The status bar at the bottom left shows "14:1 (Top Level)". Below the script pane is a console pane with tabs for Console, Terminal, and Background Jobs. The Console tab is active, showing the following R session:

```
> x <- c(1,5,2,10)
> mean(x)
[1] 4.5
> x^2
[1]  2 10  4 20
> y <- x^2
> y
[1]  2 10  4 20
>
```

Exercise

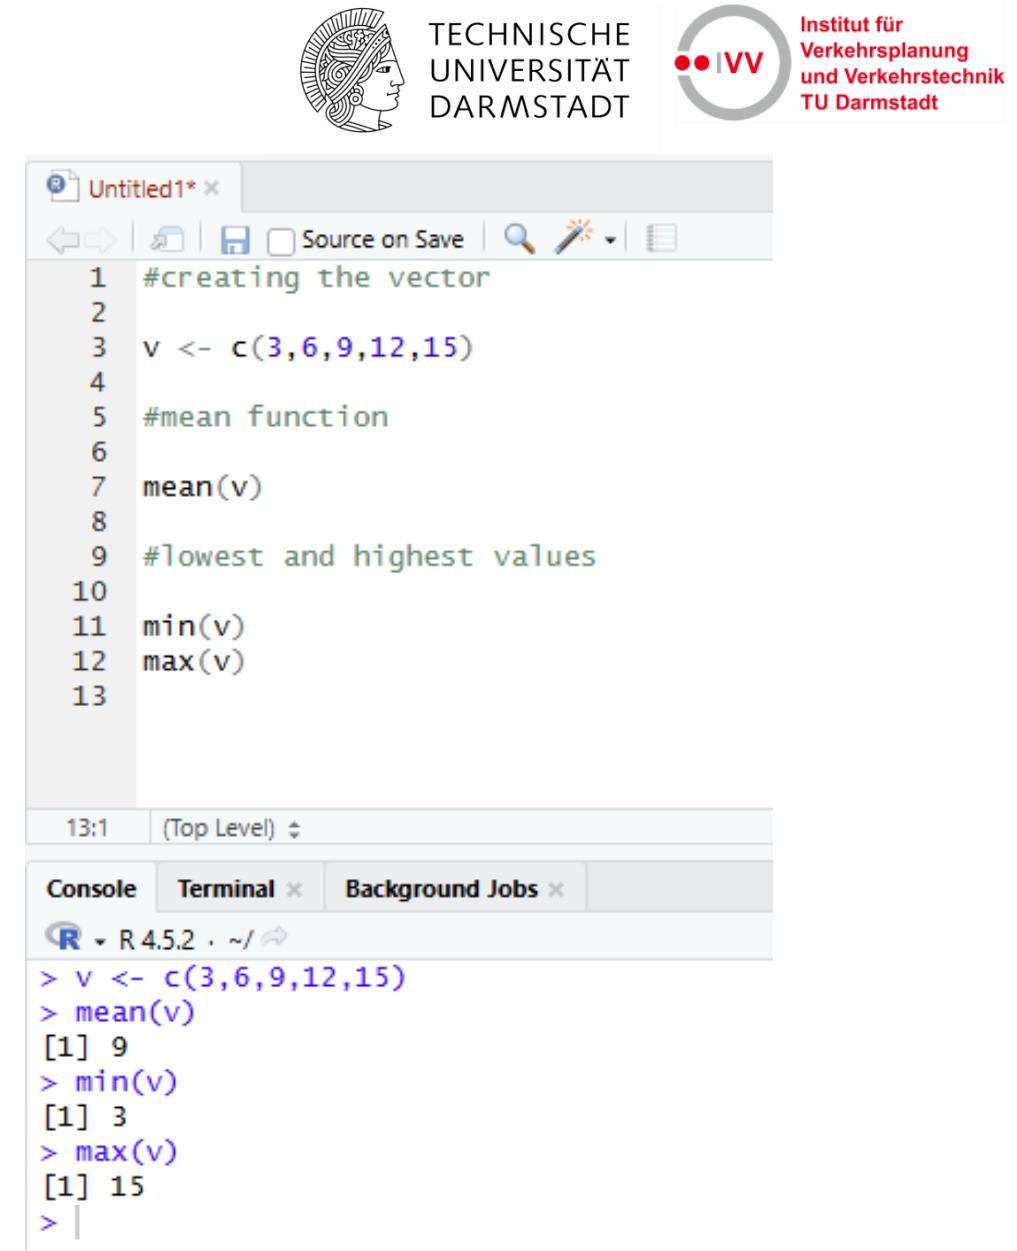
- Create a numeric vector with 5 values
- Compute mean, highest value, and lowest value
- Functions are **mean()**, **min()**, **max()**



Try this simple exercise

- Create a numeric vector with 5 values
- Compute mean, highest value, and lowest value
- Functions are `mean()`, `min()`, `max()`
- **Solution:**

```
#creating the vector v  
  
v <- c(3, 6, 9, 12, 15)  
  
#using the mean, min and max functions  
  
mean(v)  
min(v)  
max(v)
```



The screenshot shows the RStudio interface with the following details:

- Code Editor:** The "Untitled1" tab contains R code for creating a vector and computing statistics.
- Console:** The R session output shows the creation of vector `v` and the results of `mean(v)`, `min(v)`, and `max(v)`.
- Terminal:** Not used in this session.
- Background Jobs:** Not used in this session.

```
R #creating the vector  
v <- c(3, 6, 9, 12, 15)  
#mean function  
mean(v)  
#lowest and highest values  
min(v)  
max(v)  
  
> v <- c(3, 6, 9, 12, 15)  
> mean(v)  
[1] 9  
> min(v)  
[1] 3  
> max(v)  
[1] 15  
>
```

Scripts ensure Reproducibility of codes

- While working on code that you want to reuse -- always work from a script (source pane)
- **Save early, save often**
- **Comments should be used to label or describe the code**



The screenshot shows the RStudio interface. The top panel displays a script named "Untitled1" with the following R code:

```
1 #creating the vector
2
3 v <- c(3,6,9,12,15)
4
5 #mean function
6
7 mean(v)
8
9 #lowest and highest values
10
11 min(v)
12 max(v)
13
```

The bottom panel shows the "Console" tab with the following output:

```
> v <- c(3,6,9,12,15)
> mean(v)
[1] 9
> min(v)
[1] 3
> max(v)
[1] 15
>
```

What are packages?

- R comes with basic functionality
- R community develops packages
- Packages add new functions
- Package is a collection of functions
- Written for a specific goals
- We load packages when we need them

- Most packages come from CRAN (The Comprehensive R Archive Network)
- Some packages are also on GitHub
- To find the right packages, search the web with your objective..Eg:
package for creating scatterplots



Installing packages

- Install = download it using R onto your PC
- Done once for a package
- Requires internet connection
- Packages usually come from CRAN
- Packages need to **loaded** for each session
- Install the packages now: dplyr and ggplot2
 - `Install.packages("dplyr")`
 - `Install.packages("ggplot2")`
- Install ≠ load...the function for loading is **library()**
 - `Library(dplyr)`
 - `Library(ggplot2)`



Working Directory (How R finds files)



- R is always working inside one assigned folder
- This folder is called the working directory
- To check it, use the function:
 - `getwd()`
- Function to list the files (R can access) in the folder:
 - `list.files()`
- If you open an Rstudio Project, the working directory is automatically the project folder
- Otherwise it is a default folder (documents, etc.)

RStudio Project (Solution to file path problems)

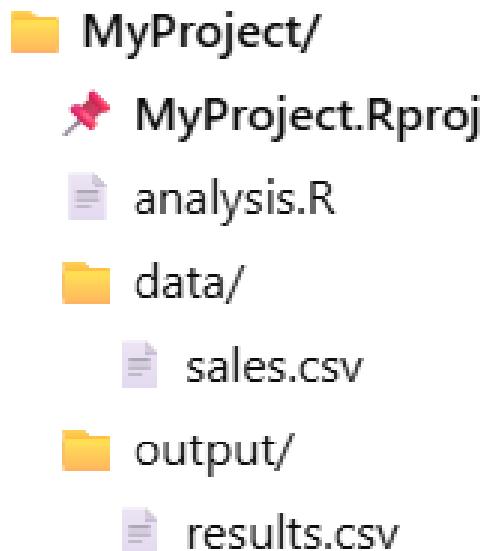
An Rstudio Project is a dedicated workspace folder.

When you use a project, Rstudio automatically:

- Sets the working directory
- Keeps scripts + data together
- Makes your work portable



Ideal Project Folder Structure



How data enters R

- Common source formats: CSV, Excel, Text files
- There are also built-in datasets (great for practice)

To import a CSV file we can use **read.csv()** function

Some functions to examine the imported data:

- **head()** function shows the first few rows of the dataset (quick preview)
- **names()** function shows the column names (what variables you have)
- **summary()** function gives a quick statistical summary of each column
- **str()** function shows the structure of the dataset – number of rows/columns + each columns **datatype** and example values



Code-along



- Lets use the built in dataset “mtcars” and try the functions to examine the data

```
data(mtcars)    # load dataset (optional)
head(mtcars)   # first rows
names(mtcars)  # column names
str(mtcars)    # structure + data types
summary(mtcars) # quick statistics
```

Output obtained
using the function:
`head(mtcars)`

```
> head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Your turn: Inspect the dataset

- Assign the dataset “iris” to a new object (hint: `df <- iris`)
- Use the data, head, names, str and summary function on this object

- How many rows and columns does the dataset have?
- What are the different variables (column names) in the dataset?



Solutions: Inspect the dataset

- Assign the dataset “iris” to a new object (hint: `df <- iris`)
- Use the `head`, `names`, `str` and `summary` function on this object

```
df <- iris

head(df)
names(df)
str(df)
summary(df)
dim(df)
```



- How many rows and columns does the dataset have? **150 rows, 5 columns**
- What are the different variables (column names) in the dataset? **Sepal.Length, Sepal.Width, Petal.Length, Petal.Width, Species**

View the data



- The function `view()` opens a dataset in a table view (similar to excel)
- Useful for quickly checking rows and columns
- It does not change your data
- Try running these codes:

```
View(df)  
View(iris)  
View(mtcars)
```

- Additional tip: use `?` To learn about a dataset or a function:

```
?mtcars  
?iris
```

Missing values in datasets : (NA)



- ‘NA’ means ‘missing value’
- When importing data, R converts blank values to NA automatically.

- Try this in RStudio:

```
x <- c(1, 2, NA, 4)  
  
mean(x)  
mean(x, na.rm = TRUE)
```

- **na.rm = TRUE** tells R to ignore missing values
- If even one value is missing, many calculations return NA unless you handle it.
Eg: mean()
- But sometimes missing values are written differently (like N/A, ?, -) and then R does not recognise them.

What is data wrangling?



- **Data wrangling = preparing data for analysis**
- Typical tasks:
 - Selecting useful columns
 - Filtering rows
 - Creating new variables
 - Fixing missing values
 - Preparing data for plots and summaries



Pipes (%>%)

The pipe `%>%` passes the **output of one step** into the **next function**. So instead of writing nested functions, we can write it step by step.

So instead of nested functions like this:

```
select(filter(mtcars, cyl == 6), mpg, hp)
```

We use this:

```
mtcars %>%  
  filter(cyl == 6) %>%  
  select(mpg, hp)
```

Read it like: take mtcars, then filter, then select



Shortcut:

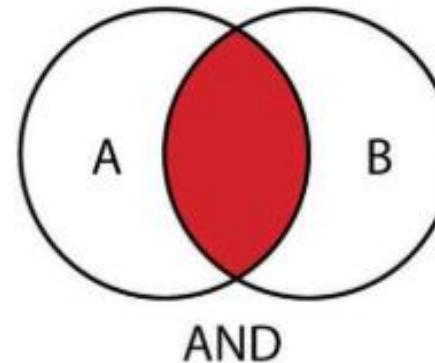
in RStudio press **Ctrl + Shift + M**
to insert `%>%`

Filtering rows with `filter()` function

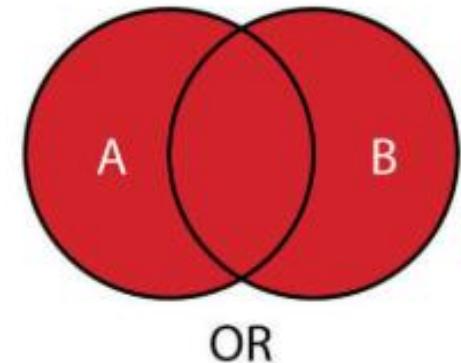
`filter()` keeps only rows that match condition(s)

Try these examples in RStudio using the “iris” dataset on flowers:

- **One condition:** keep only flowers that are “setosa”
- **AND (&) condition:** keep only flowers that are “setosa” & Sepal.Length > 5
- **OR (|) condition:** keep only flowers that are “setosa” OR have Sepal.Length > 7



AND



OR

```
library(dplyr)  
iris %>% filter(Species == "setosa")
```

```
iris %>% filter(Species == "setosa" &  
Sepal.Length > 5)
```

```
iris %>% filter(Species == "setosa" |  
Sepal.Length > 7)
```

Your turn: Filtering rows with `filter()` function

Using the built-in dataset “mtcars”, write R code to:

- **One condition:** Keep only cars with 6 cylinders (`cyl == 6`)
- **AND (&) condition:** Keep only cars with 8 cylinders (`cyl == 8`) AND fuel efficiency < 15 (`mpg < 15`)
- **OR (|) condition:** Keep cars with horsepower > 200 (`hp > 200`) OR weight > 4 (`wt > 4`)

 Hint: use `filter()` and `==`, `<`, `>`



Solutions: Filtering rows with `filter()` function

- **One condition:** Keep only cars with 6 cylinders (`cyl == 6`)
- **AND (&) condition:** Keep only cars with 8 cylinders (`cyl == 8`) AND fuel efficiency < 15 (`mpg < 15`)
- **OR (|) condition:** Keep cars with horsepower > 200 (`hp > 200`) OR weight > 4 (`wt > 4`)



```
library(dplyr)

# 1) One condition: 6 cylinders
mtcars %>% filter(cyl == 6)

# 2) AND condition: 8 cylinders AND mpg < 15
mtcars %>% filter(cyl == 8 & mpg < 15)

# 3) OR condition: hp > 200 OR wt > 4
mtcars %>% filter(hp > 200 | wt > 4)
```

Selecting columns with `select()` function

- `select()` keeps only the columns you choose (drops the rest)
- It returns a new data frame (table)
- *Try these example out:*

```
library(dplyr)

mtcars %>%
  select(mpg, hp, wt)
```

- *You may also save it as a new object:*

```
small_df <- mtcars %>% select(mpg, hp, wt)
```

Mazda RX4
Mazda RX4 Wag
Datsun 710
Hornet 4 Drive
Hornet Sportabout
Valiant
Duster 360
Merc 240D
Merc 230

mpg	cy	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
14.3	8	360.0	245	3.21	3.570	15.84	0	0	5	4
24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2

	mpg	hp	wt
Mazda RX4	21.0	110	2.620
Mazda RX4 Wag	21.0	110	2.875
Datsun 710	22.8	93	2.320
Hornet 4 Drive	21.4	110	3.215
Hornet Sportabout	18.7	175	3.440
Valiant	18.1	105	3.460
Duster 360	14.3	245	3.570
Merc 240D	24.4	62	3.190
Merc 230	22.8	95	3.150



Your turn: Selecting columns



- Use the dataset “iris”
- Select any 2 variables/columns out of the available ones
(eg: Petal.Length, Petal.Width, Species, etc.)
- Save the result into a new object called iris_small

• *Hint:*

```
library(dplyr)

iris_small <- iris %>%
  select(...)
```

• *Then check your result:*

```
head(iris_small)
```

Creating new variables with **mutate()** function

- `mutate()` creates new columns (new variables) in a dataset
- Can also modify existing columns (by overwriting them)
- *Try this example out in RStudio:*

```
library(dplyr)

iris_new <- iris %>%
  mutate(petal_ratio = Petal.Length /
Petal.Width)
```

- *Check the result:*

```
head(iris_new)
```



Original dataset:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa
9	4.4	2.9	1.4	0.2	setosa

New dataset with an added column:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species	petal_ratio
1	5.1	3.5	1.4	0.2	setosa	7.000000
2	4.9	3.0	1.4	0.2	setosa	7.000000
3	4.7	3.2	1.3	0.2	setosa	6.500000
4	4.6	3.1	1.5	0.2	setosa	7.500000
5	5.0	3.6	1.4	0.2	setosa	7.000000
6	5.4	3.9	1.7	0.4	setosa	4.250000
7	4.6	3.4	1.4	0.3	setosa	4.666667
8	5.0	3.4	1.5	0.2	setosa	7.500000
9	4.4	2.9	1.4	0.2	setosa	7.000000

Your turn: Create a new variable with **mutate()** function



- Using the dataset iris, create a new dataset called iris_new2
- Add a new column called sepal_ratio
- $\text{sepal_ratio} = \text{Sepal.Length} / \text{Sepal.Width}$
- Hint:

```
library(dplyr)

iris_new2 <- iris %>%
  mutate(sepal_ratio = ...)
```

- Then check your result

```
head(iris_new2)
```

Solution: Create a new variable with `mutate()` function



- A new column is added to the dataset
- Calculated row by row (for every row)
- *Solution:*

```
library(dplyr)

iris_new2 <- iris %>%
  mutate(sepal_ratio = Sepal.Length /
Sepal.Width)

head(iris_new2)
```

- *Optional quick check:*

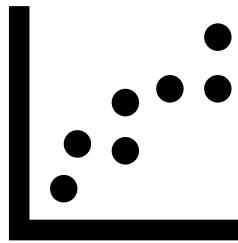
```
summary(iris_new2$sepal_ratio)
```

The **\$ operator** is used to access **one column (variable)** inside a data frame. So the code here means “Give me the sepal_ratio column from the dataset iris_new2”

Visualizations (using R)



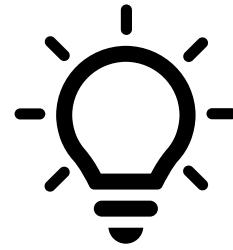
- Visualization helps us quickly see:



Patterns
(trends,
relationships)



Comparisons
(groups,
categories)



Insights
(what matters in
the data)



Outliers
(unusual values)

- We're going to use **ggplot2**, the most popular plotting package in R.

First plot in ggplot2: Scatter plot

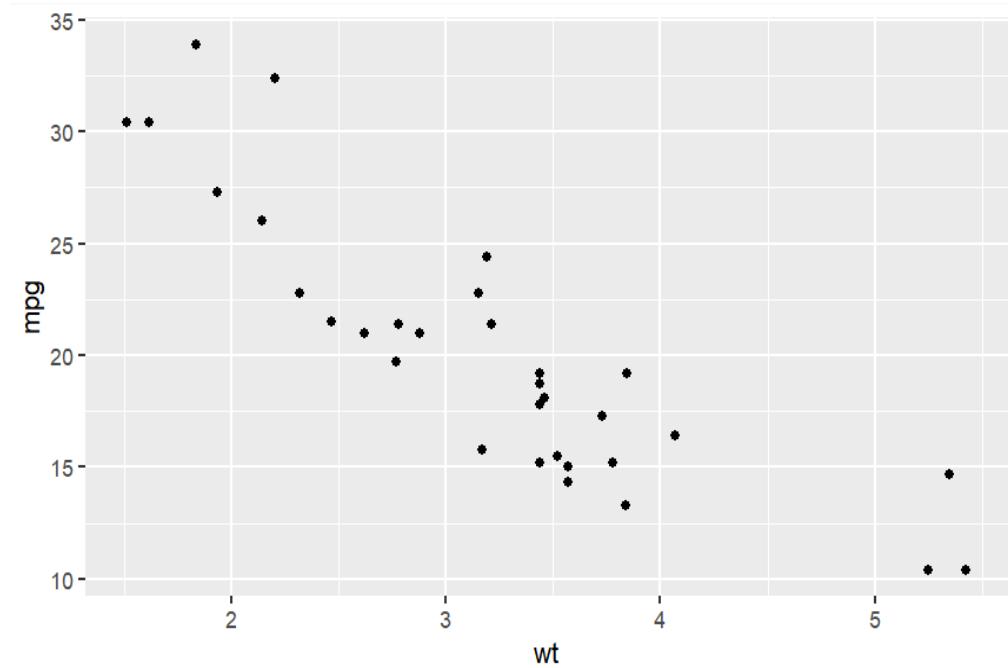
- Goal: visualize the relationship between two numeric variables
- In a scatter plot each point is one observation
- In mtcars, each point would represent a car

Ggplot2 recipe:

1. Choose dataset
2. Map variables to axes -- **aes()**
3. Choose a geometry (plot type)

```
library(ggplot2)

ggplot(mtcars, aes(x = wt, y = mpg)) +
  geom_point()
```



#5 Visualisation with ggplot2

Color encodes information (not decoration)

- Color can represent a **variable** in plots
- This helps us **compare groups** in the same plot
- We usually colour by a **category** (eg: number of cylinders)
- In mtcars, cyl (cylinders) is numeric, but we treat it as a category
- Try this code in RStudio:

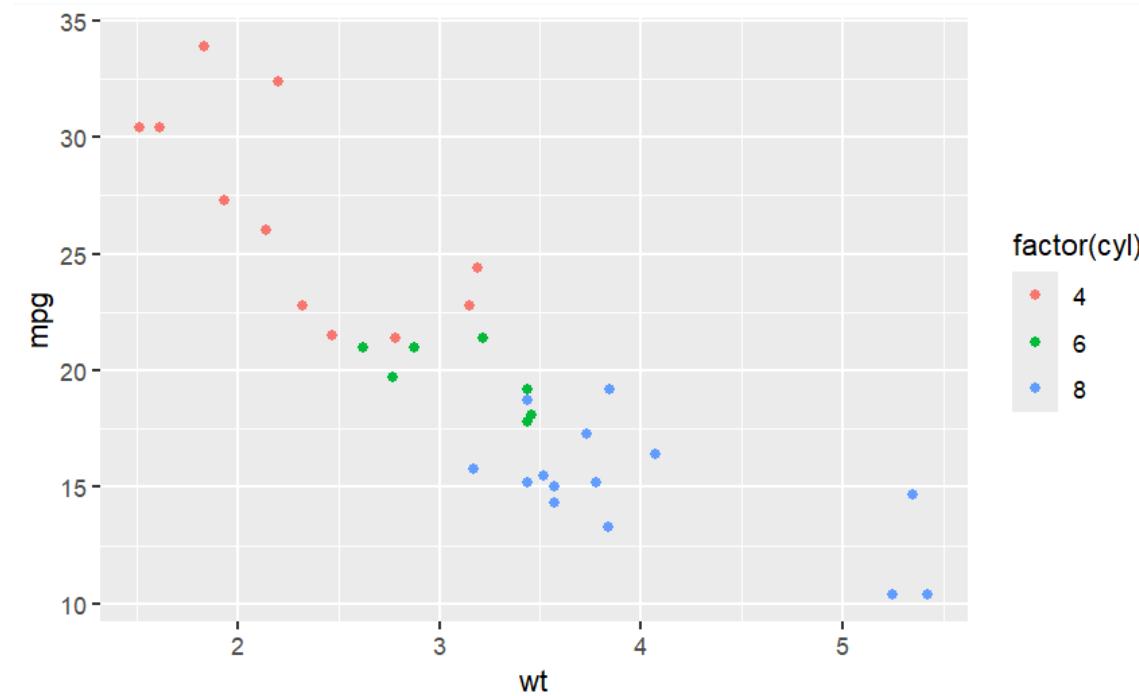
```
mtcars %>%  
  
  ggplot(aes(x = wt, y = mpg, color =  
             factor(cyl))) +  
  geom_point()
```



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt



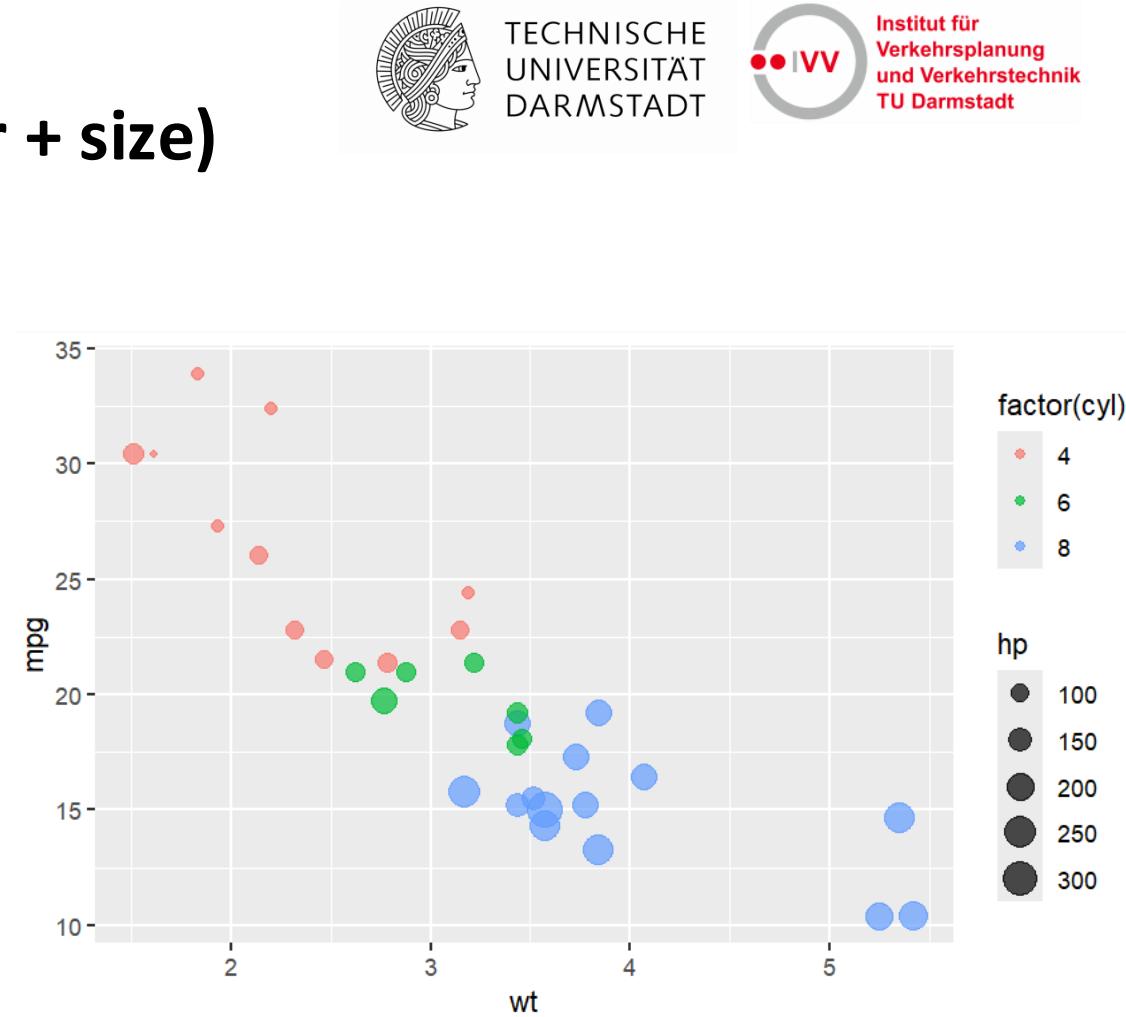
The plot now shows differences
between **4, 6, and 8 cylinder cars**

#5 Visualisation with ggplot2

More information on the same plot (colour + size)

- In addition to colour, we can use size to visualize another extra variable in a scatter plot
- Colour – categorical variable
- Size – numeric variable (eg: weight, horsepower, fuel efficiency)
- Try this code in RStudio:

```
mtcars %>%  
  ggplot(aes(x = wt, y = mpg,  
             color = factor(cyl),  
             size = hp)) +  
  geom_point(alpha = 0.7)
```



We added cylinders (colour) + horsepower (size)



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt

Your turn: Scatter plot (colour + size)

- Use mtcars to create a scatter plot
 - x-axis: wt
 - y-axis: mpg
- Add 2 extra variables
 - One for colour (categorical)
 - One for size (numeric)
- Choose options you have not used yet. E.g.:
 - colour = factor(gear) (number of gears)
 - colour = factor(am) (0 = automatic, 1 = manual)
 - size = qsec (1/4 mile time)
 - size = disp (engine displacement)



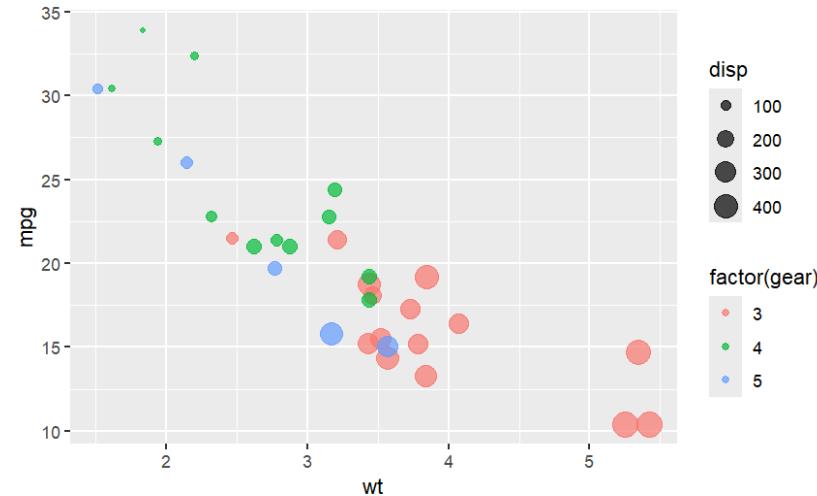
```
mtcars %>%  
  ggplot(aes(x = wt, y = mpg,  
             color = factor(gear),  
             size = disp)) +  
  geom_point(alpha = 0.7)
```

#5 Visualisation with ggplot2

Solution: Scatter plot (colour + size)

- There many correct answers
- Two example solutions:

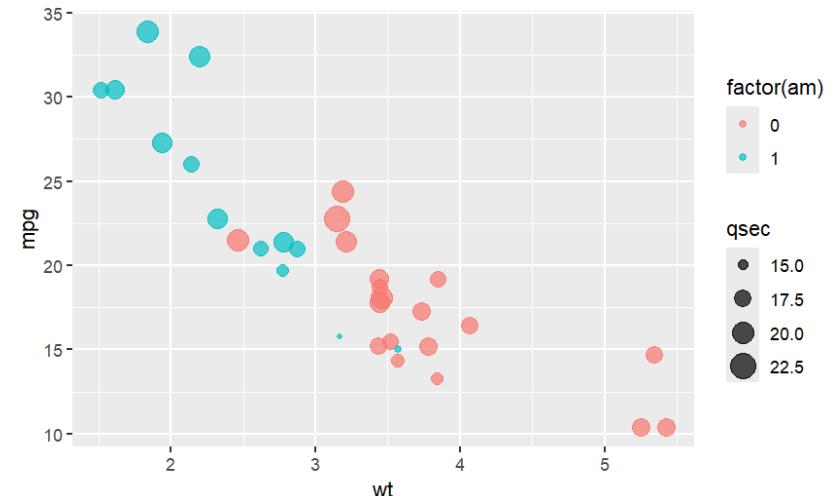
```
mtcars %>%  
  ggplot(aes(x = wt, y = mpg,  
             color = factor(gear),  
             size = disp)) +  
  geom_point(alpha = 0.7)
```



- You can swap the variables:
 - **Color:** gear or am
 - **Size:** disp or qsec

The “+” is important!!

```
mtcars %>%  
  ggplot(aes(x = wt, y = mpg,  
             color = factor(am),  
             size = qsec)) +  
  geom_point(alpha = 0.7)
```



Saving results

- We often want to save a table from R to a file
- The easiest format is CSV
- Use `write.csv()` to export a data frame
- Example (save the first 10 rows):

```
result <- mtcars %>%
  select(mpg, wt, hp) %>%
  head(10)

write.csv(result, "mtcars_sample.csv",
row.names = FALSE)
```



Saving plots

- After creating a plot, we often want to save it as an image

• **ggsave()** saves the most recent plot

• Common formats:

- PNG (best for slides + web)
- PDF (best for reports)

• Save the last plot you created using ggsave():

```
ggsave("my_plot.png", width = 6, height = 4)
```

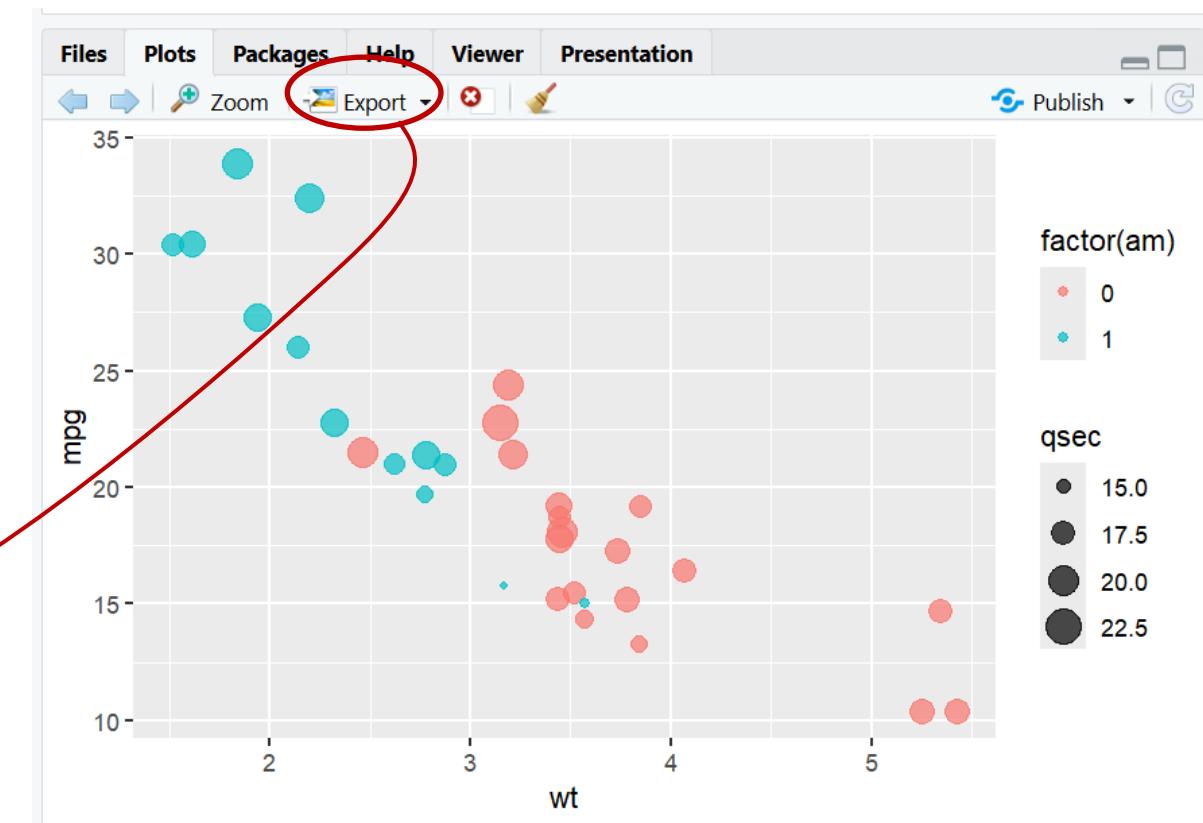
• The file is saved in your working directory

OR (manual export):

• In the **Plots** pane:

• you can click **Export**

• Then **Save as Image / Save as PDF**



When things dont work: Debugging checklist

- Read the error message
- Check if you loaded the package in use:

```
library(dplyr)  
library(ggplot2)
```

- Check if the object exists. This function lists the names of everything you created like vectors, data frames

```
ls()
```

- Inspect the data

```
head(df)  
str(df)  
names(df)
```

- Check the function help

```
?filter  
?ggplot
```

- Search the exact search message online: copy and paste the full error message into Google



Resources for learning and help

- In **RStudio** (fastest help):
 - `?Function_name` → e.g.: `?mean` → opens documentation for a function
 - `??plot` → searches help pages by keyword
- Cheat sheets (quick reference)
 - **Posit** “Data Transformation” (dplyr)
 - **Posit** “Data Visualization” (ggplot2)
- Learn by **examples and tutorials**
 - Search: “ggplot scatter plot examples”
 - Search: “dplyr filter select mutate examples”
- Community help (when you’re stuck)
 - **Stack Overflow**
 - **RStudio Community**



#6 Exporting, debugging and way forward

Final message



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt

You don't need to know everything. You need to know how to find and use what you need.



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Institut für
Verkehrsplanung
und Verkehrstechnik
TU Darmstadt



- THANK YOU
- DANKE